

```
In [173]: import cv2
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
%matplotlib inline

import os
import random
```

```
In [174]: TRAINDIR = 'C:/Users/ADMIN/Desktop/Snehil devops/train'
CATEGORIES = ["apple", "orange"]
```

```
In [175]: nrows = 150
ncolumns = 150
channels = 3
```

```
In [176]: train_data=[]
def create_training_data():
    for category in CATEGORIES:
        path = os.path.join(TRAINDIR,category)      # path to apple or orange dir
        class_num = CATEGORIES.index(category)



        for image in os.listdir(path):              # Bunch of images in the path

            image_array1 = cv2.imread(os.path.join(path,image),cv2.IMREAD_COLOR)

            new_array1 = cv2.resize(image_array1,(nrows,ncolumns))      # Normal


            color_img = cv2.cvtColor(new_array1, cv2.COLOR_BGR2RGB)
            imgplot=plt.imshow(color_img)

            train_data.append([color_img,class_num])      # Collecting all train
            plt.show()
create_training_data()
```




```
In [177]: random.shuffle(train_data)
```

```
In [178]: plt.imshow(train_data[1][0])  
plt.show()
```



```
In [179]: plt.imshow(train_data[2][0])  
plt.show()
```



```
In [180]: X = []  
y = []  
  
for feature, label in train_data:  
    X.append(feature)  
    y.append(label)  
  
X = np.array(X).reshape(-1, n_rows, n_columns, channels)  
y = np.array(y)
```

In [181]: X[0]

```
Out[181]: array([[255, 255, 255],
                 [255, 255, 255],
                 [255, 255, 255],
                 ...,
                 [255, 255, 255],
                 [255, 255, 255],
                 [255, 255, 255]],

                [[255, 255, 255],
                 [255, 255, 255],
                 [255, 255, 255],
                 ...,
                 [255, 255, 255],
                 [255, 255, 255],
                 [255, 255, 255]],

                [[255, 255, 255],
                 [255, 255, 255],
                 [255, 255, 255],
                 ...,
                 [255, 255, 255],
                 [255, 255, 255],
                 [255, 255, 255]],

                ...,

                [[255, 255, 255],
                 [255, 255, 255],
                 [255, 255, 255],
                 ...,
                 [255, 255, 255],
                 [255, 255, 255],
                 [255, 255, 255]],


                [[255, 255, 255],
                 [255, 255, 255],
                 [255, 255, 255],
                 ...,
                 [255, 255, 255],
                 [255, 255, 255],
                 [255, 255, 255]],

                [[255, 255, 255],
                 [255, 255, 255],
                 [255, 255, 255],
                 ...,
                 [255, 255, 255],
                 [255, 255, 255],
                 [255, 255, 255]]], dtype=uint8)
```

In [182]: y

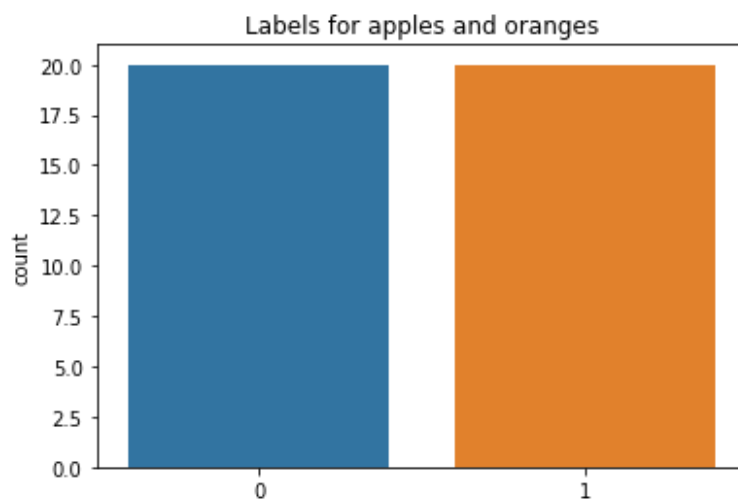
```
Out[182]: array([1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1,
                 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0])
```

```
In [183]: plt.figure(figsize=(20,20))
columns = 5
for i in range(columns):
    plt.subplot(5/columns + 1, columns, i + 1)
    plt.imshow(X[i])
```



```
In [184]: import seaborn as sns
sns.countplot(y)
plt.title('Labels for apples and oranges')
```

Out[184]: Text(0.5, 1.0, 'Labels for apples and oranges')



```
In [185]: X.shape
```

Out[185]: (40, 150, 150, 3)

```
In [186]: y.shape
```

Out[186]: (40,)

```
In [187]: from sklearn.model_selection import train_test_split
X_train, X_val, y_train, y_val = train_test_split(X,y,test_size=0.2,random_state=
```

```
In [188]: print(X_train.shape)
          print(X_val.shape)
          print(y_train.shape)
          print(y_val.shape)
```

```
(32, 150, 150, 3)
(8, 150, 150, 3)
(32,)
(8,)
```

```
In [189]: #get the length of the train and validation data
          ntrain = len(X_train)
          nval = len(X_val)

          #We will use a batch size of 32. Note: batch size should be a factor of 2.**4,8,
          batch_size = 32

          print(ntrain)
          print(nval)
```

```
32
8
```

```
In [190]: from keras import layers
          from keras import models
          from keras import optimizers
          from keras.preprocessing.image import ImageDataGenerator
          from keras.preprocessing.image import img_to_array, load_img

          model = models.Sequential()
          model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)))
          model.add(layers.MaxPooling2D((2, 2)))
          model.add(layers.Conv2D(64, (3, 3), activation='relu'))
          model.add(layers.MaxPooling2D((2, 2)))
          model.add(layers.Conv2D(128, (3, 3), activation='relu'))
          model.add(layers.MaxPooling2D((2, 2)))
          model.add(layers.Conv2D(128, (3, 3), activation='relu'))
          model.add(layers.MaxPooling2D((2, 2)))
          model.add(layers.Flatten())
          model.add(layers.Dropout(0.5)) #Dropout for regularization
          model.add(layers.Dense(512, activation='relu'))
          model.add(layers.Dense(1, activation='sigmoid')) #Sigmoid function at the end be
```

In [191]: `model.summary()`

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
=====		
conv2d_21 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_21 (MaxPooling)	(None, 74, 74, 32)	0
conv2d_22 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_22 (MaxPooling)	(None, 36, 36, 64)	0
conv2d_23 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_23 (MaxPooling)	(None, 17, 17, 128)	0
conv2d_24 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_24 (MaxPooling)	(None, 7, 7, 128)	0
flatten_6 (Flatten)	(None, 6272)	0
dropout_6 (Dropout)	(None, 6272)	0
dense_11 (Dense)	(None, 512)	3211776
dense_12 (Dense)	(None, 1)	513
=====		
Total params: 3,453,121		
Trainable params: 3,453,121		
Non-trainable params: 0		

In [192]: `model.compile(loss='binary_crossentropy', optimizer=optimizers.RMSprop(lr=1e-4),`

In [193]: `train_datagen = ImageDataGenerator(rescale=1/255,  
rotation_range=40,  
width_shift_range=0.2,  
height_shift_range=0.2,  
shear_range=0.2,  
zoom_range=0.2,  
horizontal_flip=True,)  
  
val_datagen = ImageDataGenerator(rescale=1/255)`

In [194]: `train_generator = train_datagen.flow(X_train, y_train, batch_size=batch_size)  
val_generator = val_datagen.flow(X_val, y_val, batch_size=batch_size)`

```
In [195]: history = model.fit_generator(train_generator,
                                         steps_per_epoch=ntrain // batch_size,
                                         epochs=64,
                                         validation_data=val_generator,
                                         validation_steps=nval // batch_size)

1/1 [=====] - 1s 1s/step - loss: 0.0644 - acc: 1.0000
0 - val_loss: 0.0842 - val_acc: 1.0000
Epoch 37/64
1/1 [=====] - 1s 1s/step - loss: 0.0914 - acc: 1.0000
0 - val_loss: 0.0863 - val_acc: 1.0000
Epoch 38/64
1/1 [=====] - 1s 1s/step - loss: 0.1949 - acc: 0.906
2 - val_loss: 0.3317 - val_acc: 0.8750
Epoch 39/64
1/1 [=====] - 1s 1s/step - loss: 0.1939 - acc: 0.875
0 - val_loss: 0.1744 - val_acc: 1.0000
Epoch 40/64
1/1 [=====] - 1s 1s/step - loss: 0.2132 - acc: 0.906
2 - val_loss: 0.0537 - val_acc: 1.0000
Epoch 41/64
1/1 [=====] - 1s 1s/step - loss: 0.0620 - acc: 1.0000
0 - val_loss: 0.0253 - val_acc: 1.0000
Epoch 42/64
1/1 [=====] - 1s 1s/step - loss: 0.0634 - acc: 0.968
8 - val_loss: 0.0270 - val_acc: 1.0000
```


```
In [196]: #lets plot the train and val curve
#get the details form the history object
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(1, len(acc) + 1)

#Train and validation accuracy
plt.plot(epochs, acc, 'b', label='Training accuracy')
plt.plot(epochs, val_acc, 'r', label='Validation accuracy')
plt.title('Training and Validation accuracy')
plt.legend()

plt.figure()
#Train and validation loss
plt.plot(epochs, loss, 'b', label='Training loss')
plt.plot(epochs, val_loss, 'r', label='Validation loss')
plt.title('Training and Validation loss')
plt.legend()

plt.show()
```



```
In [214]: TRAINDIR2 = 'C:/Users/ADMIN/Desktop/Snehil devops/train2'
CATEGORIES = ["apple", "orange"]
```



```
In [215]: nrows = 150
          ncolumns = 150
          channels = 3
```

```
In [216]: train_data=[]
          def create_training_data():
              for category in CATEGORIES:
                  path = os.path.join(TRAINDIR2,category)           # path to apple or orange dir
                  class_num = CATEGORIES.index(category)

                  for image in os.listdir(path):                   # Bunch of images in the path

                      image_array1 = cv2.imread(os.path.join(path,image))

                      new_array1 = cv2.resize(image_array1,(nrows,ncolumns))           # Normalizing the image

                      train_data.append([new_array1,class_num])           # Collecting all training data

          create_training_data()
```

```
In [217]: random.shuffle(train_data)
```

```
In [218]: X = []
          y = []

          for feature,label in train_data:
              X.append(feature)
              y.append(label)

          X = np.array(X).reshape(-1,nrows,ncolumns,channels)
          y = np.array(y)
```

```
In [219]: from sklearn.model_selection import train_test_split
          X_train, X_val, y_train, y_val = train_test_split(X,y,test_size=0.2,random_state=42)
```

```
In [220]: from keras import layers
          from keras import models
          from keras import optimizers
          from keras.preprocessing.image import ImageDataGenerator
          from keras.preprocessing.image import img_to_array, load_img

          model = models.Sequential()
          model.add(layers.Conv2D(32, (3, 3), activation='relu',input_shape=(150, 150, 3)))
          model.add(layers.MaxPooling2D((2, 2)))
          model.add(layers.Conv2D(64, (3, 3), activation='relu'))
          model.add(layers.MaxPooling2D((2, 2)))
          model.add(layers.Conv2D(128, (3, 3), activation='relu'))
          model.add(layers.MaxPooling2D((2, 2)))
          model.add(layers.Conv2D(128, (3, 3), activation='relu'))
          model.add(layers.MaxPooling2D((2, 2)))
          model.add(layers.Flatten())
          model.add(layers.Dropout(0.5))           #Dropout for regularization
          model.add(layers.Dense(512, activation='relu'))
          model.add(layers.Dense(1, activation='sigmoid'))           #Sigmoid function at the end because we have only two classes
```

```
In [221]: model.compile(loss='binary_crossentropy', optimizer=optimizers.RMSprop(lr=1e-4),
```

```
In [222]: train_datagen = ImageDataGenerator(rescale=1/255,
                                             rotation_range=40,
                                             width_shift_range=0.2,
                                             height_shift_range=0.2,
                                             shear_range=0.2,
                                             zoom_range=0.2,
                                             horizontal_flip=True,)

val_datagen = ImageDataGenerator(rescale=1/255)
```

```
In [223]: train_generator = train_datagen.flow(X_train, y_train, batch_size=batch_size)
val_generator = val_datagen.flow(X_val, y_val, batch_size=batch_size)
```

```
In [224]: history = model.fit_generator(train_generator,
                                       steps_per_epoch=ntrain // batch_size,
                                       epochs=64,
                                       validation_data=val_generator,
                                       validation_steps=nval // batch_size)
```


```
0000 - val_loss: 0.3210 - val_acc: 1.0000
Epoch 37/64
1/1 [=====] - 0s 228ms/step - loss: 0.3466 - acc: 1.
0000 - val_loss: 0.6908 - val_acc: 1.0000
Epoch 38/64
1/1 [=====] - 0s 249ms/step - loss: 0.3102 - acc: 1.
0000 - val_loss: 0.4033 - val_acc: 1.0000
Epoch 39/64
1/1 [=====] - 0s 237ms/step - loss: 0.2068 - acc: 1.
0000 - val_loss: 0.4103 - val_acc: 1.0000
Epoch 40/64
1/1 [=====] - 0s 217ms/step - loss: 0.2023 - acc: 1.
0000 - val_loss: 0.3812 - val_acc: 1.0000
Epoch 41/64
1/1 [=====] - 0s 237ms/step - loss: 0.1621 - acc: 1.
0000 - val_loss: 0.5058 - val_acc: 1.0000
Epoch 42/64
1/1 [=====] - 0s 241ms/step - loss: 0.2003 - acc: 1.
0000 - val_loss: 0.2233 - val_acc: 1.0000
Epoch 43/64
```

```
In [225]: test_datagen=ImageDataGenerator(rescale=1./255)
```

```

In [226]: i = 0
text_labels = []
plt.figure(figsize=(30,20))
for batch in test_datagen.flow(X, batch_size=1):
    pred = model.predict(batch)
    if pred > 0.5:
        text_labels.append('orange')
    else:
        text_labels.append('apple')
    plt.subplot(5 / columns + 1, columns, i + 1)
    plt.title('This is a ' + text_labels[i])
    color_img = cv2.cvtColor(batch[0], cv2.COLOR_BGR2RGB)
    imgplot=plt.imshow(color_img)
    i += 1
    if i % 10 == 0:
        break
plt.show()

```



In [ ]:

